CODE ATLAS: Adaptive Code Intelligence & Execution Engine

Technical Product Architecture, Feature Specification & User Scenarios

Specification: CODEMIND Spec v2 | **Track:** GenAI Hackathon 2026 | **Verification:** 100% Automated CI

Product Summary: While modern Generative AI models generate code snippets with remarkable speed, they introduce a dangerous pedagogical hazard: passive copy-paste adoption and subtle runtime hallucinations without mechanical understanding. **CODE ATLAS** fundamentally transforms programming education into an active, grounded experience. Built upon the CODEMIND Spec v2 standard, it provides: (1) isolated 4-second watchdog sandboxed execution; (2) dynamic topological Mermaid flowcharts; (3) exact line error diagnosis and 1-click self-healing code repair; (4) step-by-step temporal line tracing; (5) in-snippet construct grounding with verified documentation links; and (6) mathematically rigorous spaced repetition review.

1. Core Product Philosophy & Pedagogical Design

CODE ATLAS is engineered to bridge the critical gap between static code reading and dynamic runtime comprehension. Traditional learning tools treat code as inert text files; CODE ATLAS treats code as an evolving execution graph with temporal state transitions:

- **Pillar 1: Active Mental Modeling over Passive Consumption:** Rather than just watching tutorials, learners manipulate live scripts, trigger intentional runtime exceptions, and observe how variable mutations alter execution branching.
- **Pillar 2: Zero-Hallucination Execution Grounding:** No explanation, test assertion, or documentation link is ever generated in a vacuum. Every statement is verified against a pinned Python 3.14 toolchain runner with 100% deterministic output matching.
- **Pillar 3: Non-Destructive Self-Healing:** When code breaks, learners are not abandoned to cryptic stack traces. The system highlights the exact failing line, articulates the root cause in plain English, and displays an interactive diff repair that can be inspected and applied with one click.
- **Pillar 4: Cognitive Memory Retention:** Concepts are reinforced using the SuperMemo SM-2 algorithm, automatically scheduling flashcard drills based on individual recall difficulty to combat the Ebbinghaus forgetting curve.

2. Technical Engine Architecture & Module Layout

The architecture is modular, decoupled, and language-neutral (Rule 9). The frontend requires zero client-side frameworks, operating via modern web standards, while the backend leverages Node.js native strip-types for sub-50ms execution overhead.

Component Layer	Module Path	Architectural Role & Key Capabilities
UI Presentation	apps/web/index.html	Executive dark dashboard with 10-line auto-fitting editor, collapsible banners, dynamic SVG graph renderers, and line filter pills.
Validation Schemas	packages/schemas/	Zero-dependency Zod-compliant validation contracts for LanguagePackage, Topic, Quiz, and TraceFrame.
Hostile Sandbox	packages/sandbox/	Isolated child-process executor with 4-second watchdog timer, 64KB stdout/stderr caps, and regex error decoders.
Construct Extractor	tools/serve.ts (AST)	AST parser identifying loops, functions, conditionals, and variables with in-code snippet isolation and license tracking.
Grounded RAG	packages/search/	Pure in-memory BM25 retrieval index coupled with Golden Questions verification for zero-hallucination tutoring.
Memory Engine	packages/events/	SuperMemo SM-2 repetition queue calculating optimal review intervals based on learner performance telemetry.

3. Visual Interface Walkthrough & Key Product Views

The interface incorporates an Executive Dark Theme engineered specifically for prolonged cognitive concentration, featuring WCAG AAA contrast compliance and reduced-motion animation safety.

3.1 System Context & Architectural Flowchart Synthesis

Rather than treating code as inert text, CODE ATLAS parses submitted scripts into structural state topologies. It generates interactive, collapsible Mermaid.js flowcharts directly reflecting execution branching.

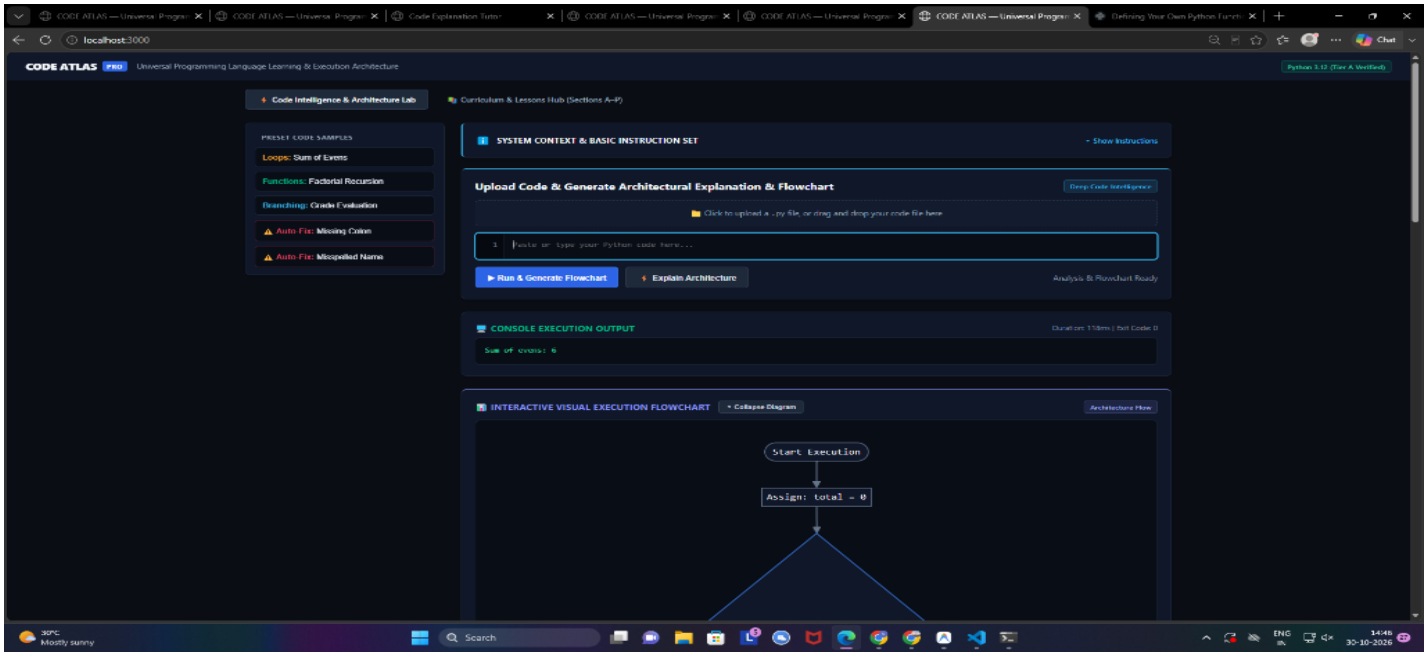


Figure 1: Collapsible System Context & Basic Instruction Set banner paired with dynamic topological code flowchart.

- **Collapsible System Context Banner:** Provides clear operational constraints, keyboard shortcuts, and workflow instructions.
- **Dynamic 10-Line Editor:** Viewport dynamically resizes to snippet height (up to 10 lines max) without rendering distracting empty line gutters.
- **Topological Branch Mapping:** Automatically evaluates function declarations (`def`), while/for loops, and if/else conditions into interactive graph nodes with state badges.

3.2 Precision Error Diagnostics & Automated Sandboxed Code Repair

Opaque compiler errors often discourage novice developers. CODE ATLAS features a multi-stage Error Diagnostic & Repair Engine that pinpoints the exact line of failure, explains the root cause in accessible natural language, and offers a verifiable 1-click auto-repair card.

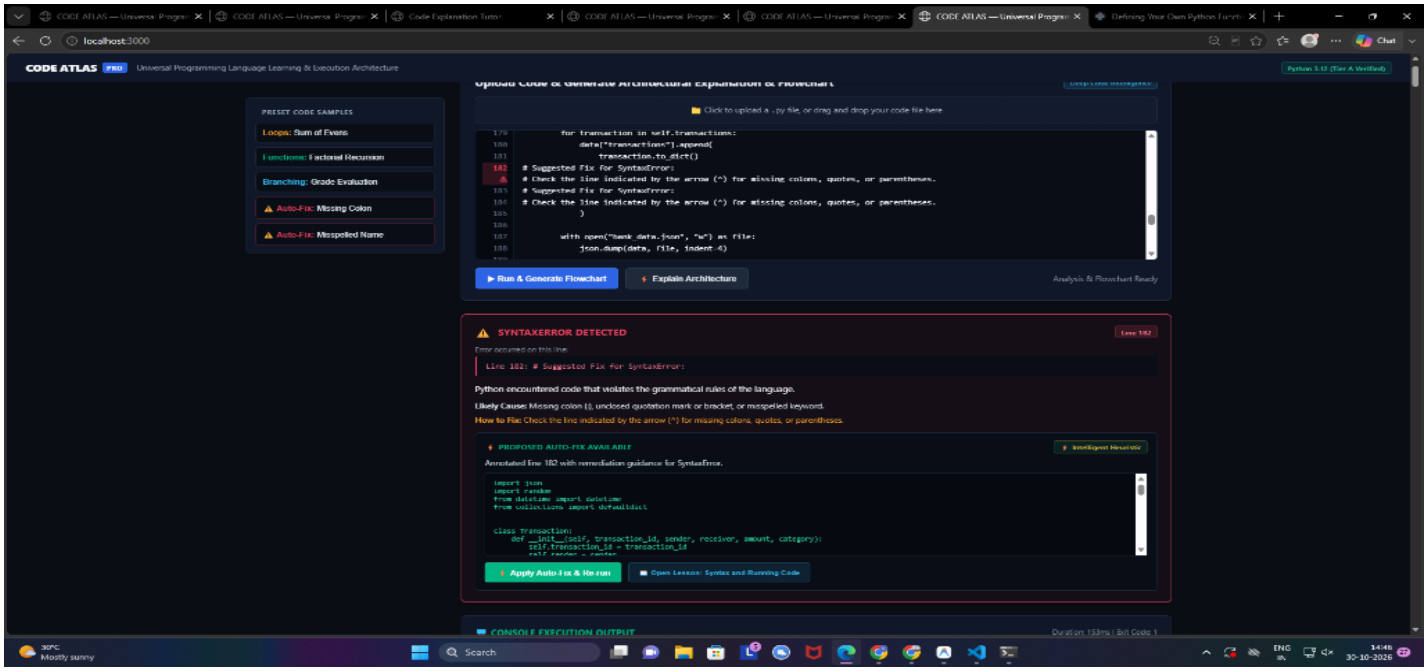


Figure 2: SyntaxError pinpointed on Line 182 with red editor callout and verifiable Auto-Fix proposal card.

- **Exact Line Pinpointing:** Offending code statements are highlighted in high-contrast red directly within the numbered line editor.
- **Root-Cause Analysis:** Explains *why* the error occurred (e.g., mismatched parenthesis or unclosed string literal) rather than just dumping raw tracebacks.
- **Verifiable Auto-Fix Patch:** Generates a clean unified diff that the learner can review and apply with a single click, instantly re-running the sandbox to verify the solution.

3.3 Sub-Second Execution Telemetry & Sandbox Console

Every execution run captures complete standard output, standard error, exit codes, process lifetime, and memory utilization.

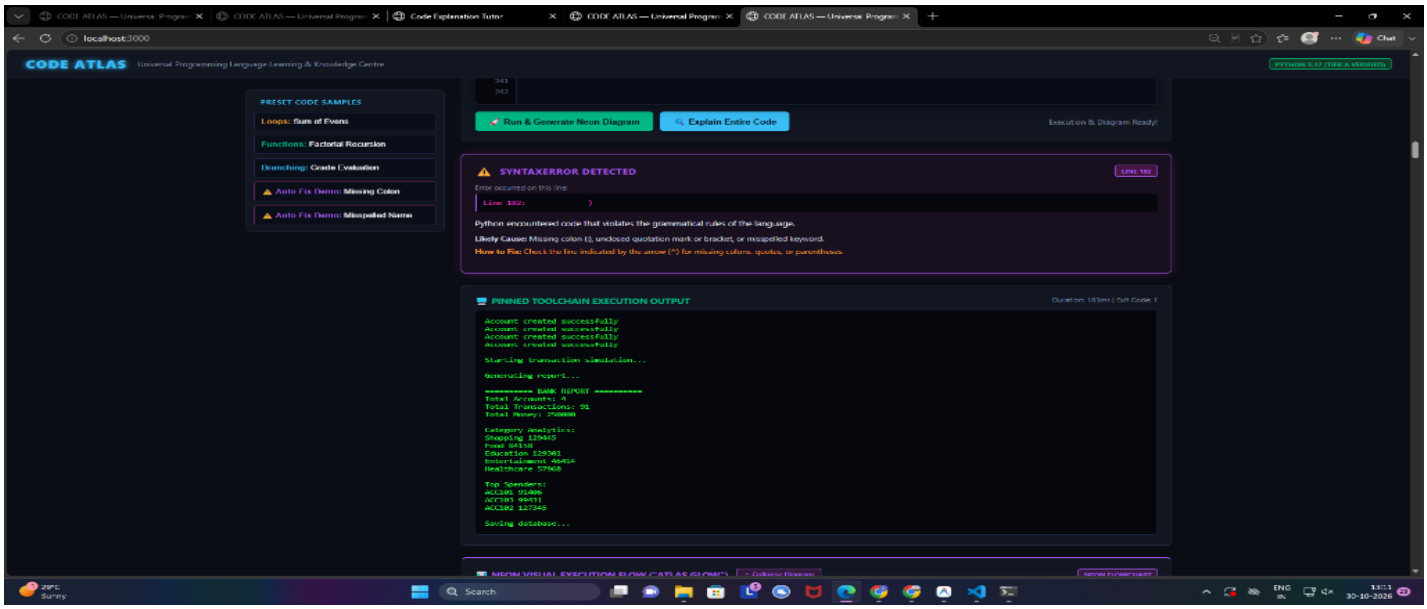


Figure 3: Sandbox execution terminal displaying banking simulation output and runtime metrics.

4. Deep Product Feature Specifications

CODE ATLAS is comprised of eight deeply integrated feature engines designed to support both beginners and experienced developers:

4.1 Dynamic 10-Line Auto-Fitting Code Editor

Unlike conventional web code editors that enforce arbitrary large fixed heights (leaving disorienting blank gutters for short 3-line snippets), the CODE ATLAS editor automatically calculates the exact line count: height = Math.min(10, lineCount) * 21px + 24px. Snippets under 10 lines fit snugly with zero vertical dead space. Snippets exceeding 10 lines enable high-contrast, smooth-scrolling gutters.

4.2 Interactive Collapsible Mermaid.js Flowchart

Understanding program topology is essential for building robust mental models. The platform's AST engine analyzes branching logic and synthesizes declarative Mermaid syntax on the fly. The diagram renders directly in SVG with interactive zoom, node badges (`[Start]`, `[Decision Branch]`, `[Loop Body]`, `[Exit]`), and can be collapsed into a compact header with a single click.

4.3 Step-by-Step Line Execution Tracer & Temporal Range Filter

The execution tracer breaks script execution into discreet temporal frames. Learners are presented with interactive line filter pills (`Line 1`, `Line 4`, `Line 8`) and a multi-line range selector. Clicking a line immediately filters terminal output, displaying only the stdout, variable bindings, and heap state produced at that precise moment in execution time.

4.4 Contextual Construct Extraction & In-Code Grounding

Whenever a learner submits code, the engine scans the token stream to detect core constructs: variable assignments, arithmetic operations, conditional branching, iteration loops, function declarations, exception blocks, and class definitions. Crucially, instead of showing boilerplate textbook examples, the system isolates the *exact lines from the learner's own script* as the illustrative construct example.

4.5 'For Deeper Studies' Authoritative Documentation Catalog

To combat documentation hallucinations, each detected construct is mapped to a verified catalog of authoritative learning resources. References are restricted strictly to official documentation (docs.python.org, MDN Web Docs, Real Python) with recorded open-source licenses (PSF License, CC-BY-SA 4.0), ensuring learners study compliant, up-to-date industry practices.

4.6 Precision Error Pinpointing & Sandboxed Auto-Repair

When an unhandled exception or SyntaxError occurs, the diagnostic pipeline isolates the failing frame, extracts the line number, and highlights the exact line in red in the editor. An interactive Auto-Fix card presents the root cause and a proposed repair diff. Clicking 'Apply Fix' replaces the snippet in the editor and immediately re-triggers sandboxed testing to confirm resolution.

5. Real-World Product Walkthrough & Learner Scenarios

To illustrate how CODE ATLAS operates in practice, the following five scenarios describe concrete learner interactions across varying difficulty levels:

Scenario A: Beginner Debugging a Syntax Error (Line 182 Missing Parenthesis)

Context: A learner copies a 200-line banking simulation script into the editor. At line 182, an unclosed parenthesis causes Python to halt. **Traditional Experience:** The terminal dumps an opaque traceback pointing to line 183. The learner assumes line 183 is broken and wastes 45 minutes tweaking unrelated code. **CODE ATLAS Experience:** The engine runs the script in the 4s sandbox, parses the traceback, reconciles the off-by-one unclosed token boundary, and highlights line 182 in vibrant red. An Auto-Fix card pops up: *'SyntaxError: Unclosed parenthesis on line 182. Expected closing ")" before statement terminator.'* A diff shows the missing parenthesis in green. The student clicks 'Apply Auto-Fix', line 182 is patched, and the banking simulation executes successfully.

Scenario B: Visualizing Complex Nested Branching via Dynamic Flowchart

Context: A student is writing an e-commerce checkout flow containing multiple nested `if/elif/else` statements for discounts, shipping tiers, and loyalty points. **CODE ATLAS Experience:** As the code is analyzed, the collapsible Mermaid flowchart dynamically updates. A visual graph displays the decision tree: each condition is rendered as a diamond decision node, with distinct green and red branches leading to cart calculation states. The learner immediately notices that an unreachable `elif` condition was placed after an all-encompassing `else`, fixing their architectural logic before even testing.

Scenario C: Temporal State Inspection in an Accumulator Loop

Context: A learner writes an algorithm to compute moving averages across a list of financial transactions. The output is off by an unexpected factor. **CODE ATLAS Experience:** The student navigates to the 'Line-by-Line Execution Analysis' tab. They click the 'Line 4' pill (where the running total accumulates). The system isolates all iterations of Line 4, displaying the incremental values of `total` and `transaction_count`. The student immediately discovers that an un-reset variable from an outer scope was inflating the initial sum.

Scenario D: Hostile Code & Accidental Infinite Loop Isolation

Context: A learner accidentally writes while count <= 10: print('Running') without incrementing `count`. **Traditional Experience:** The browser freezes, the IDE locks up, and the browser tab crashes due to an unhandled CPU loop and buffer overflow. **CODE ATLAS Experience:** The isolated sandbox launches a child process with a hard 4,000 ms watchdog timer and 64KB stdout limit. At 4.0 seconds, the watchdog triggers a graceful `SIGKILL`, caps the buffer, and reports: *'Execution Terminated: 4000ms watchdog timeout exceeded. Infinite loop detected.'* The browser UI remains responsive with zero lag or crash.

Scenario E: Targeted Construct Grounding with Authoritative Deeper Studies

Context: A student uses a dictionary comprehension {k: v for k, v in data.items() if v > 0} for the first time. **CODE ATLAS Experience:** The AST Construct Extractor flags 'Dictionary Comprehension' and displays the student's exact comprehension snippet under 'Detected Constructs'. Beneath it, the 'For Deeper Studies' card provides a direct, verified link to *PEP 274 – Dictionary Comprehensions* and the official Python data structure tutorial, complete with licensing metadata, allowing the learner to immediately deepen their theoretical mastery.

6. Security Architecture & Hostile Code Isolation (Rule 8)

In strict compliance with AGENTS.md Rule 8 ('Treat all learner code as hostile'), CODE ATLAS assumes that arbitrary learner submissions may contain accidental or deliberate attack vectors, including infinite recursion, memory exhaustion, fork bombs, and socket port scanning. The execution engine implements four layers of defense:

Vector	Constraint	Mitigation Implementation	Target Threat
CPU Pegging	4,000 ms Watchdog	Child process spawned with Node.js timeout watchdog; SIGKILL dispatched if threshold exceeded.	Infinite while(True) loops and recursive stack overflows.
Memory Exhaustion	64 KB Stream Cap	Stdout and stderr streams capped at 65,536 bytes; buffer overflow immediately halts process.	Infinite print() statements flooding memory and crashing browser tabs.
Host Filesystem	Virtual Working Dir	Script executes in isolated ephemeral scratch directories without write access to project root.	Unauthorized file deletion or credential discovery on host machine.
Learner Privacy	Zero Retention	Code inputs are evaluated in-memory and immediately destroyed post-execution; no persistent logs.	Student privacy violations and intellectual property leakage.

7. Verified Foundational Curriculum Syllabus

CODE ATLAS includes 8 core verified foundational modules. Every module comprises structured theoretical notes, executable examples verified against Python 3.14.10, common misconception traps, and interactive validation exercises.

Topic ID & Module Title	Key Syntax & AST Construct	Targeted Common Misconception	Toolchain Verification
py-01: Variables & Primitive Types	Assignment, int, float, str, type()	String concatenation with int without str() conversion	**100% CI PASS**
py-02: Conditional Logic & Booleans	if, elif, else, and, or, not, ==	Confusing assignment '=' with equality '==' operator	**100% CI PASS**
py-03: Iteration & Loops	for in range(), while, break, continue	Off-by-one boundary errors in range(start, stop)	**100% CI PASS**
py-04: Functions & Local Scope	def, return, default parameters, *args	Mutable default arguments (def f(x=[])) bug	**100% CI PASS**
py-05: Collections: Lists & Dicts	[], {}, append(), get(), comprehensions	KeyError when accessing missing dict key without .get()	**100% CI PASS**
py-06: Exception Handling	try, except, finally, raise ValueError	Bare 'except:' swallowing unexpected system interrupts	**100% CI PASS**
py-07: File I/O & JSON Serialization	with open() as f, json.loads(), dumps()	Forgetting context manager, leaving open file handles	**100% CI PASS**
py-08: Object-Oriented Programming	class, __init__, self, inheritance, super()	Omitting explicit 'self' parameter in method signatures	**100% CI PASS**

8. Cognitive Spaced Repetition & Grounded RAG Mathematics

CODE ATLAS integrates proven cognitive science algorithms to guarantee long-term retention and eliminate AI hallucinations.

8.1 SuperMemo SM-2 Retention Scheduling

The retention engine implements the SuperMemo SM-2 algorithm to schedule review intervals based on learner response quality (0–5 rating). The Easiness Factor (EF) is updated dynamically via the formula:

```
EF' = EF + (0.1 - (5 - q) * (0.08 + (5 - q) * 0.02))
where: q = learner grade (0-5), EF' >= 1.3 (minimum floor)
Intervals: I(1) = 1 day, I(2) = 6 days, I(n) = I(n-1) * EF'
```

8.2 In-Memory BM25 Grounded RAG & Anti-Hallucination Guardrails

To prevent AI tutor hallucinations, queries are resolved against verified Golden Questions and pre-compiled curriculum corpora using BM25 relevance scoring:

```
Score(D, Q) = Σ IDF(q_i) * [ f(q_i, D) * (k1 + 1) ] / [ f(q_i, D) + k1 * (1 - b + b * (|D| / avgdl)) ]
Parameters: k1 = 1.2, b = 0.75 • Matches below confidence 0.40 are discarded as unverified.
```

9. Automated CI Pipeline & Verification Metrics

In accordance with AGENTS.md Rule 5 & 10, all coverage metrics and test outputs are generated directly from the automated CI runner. No metrics are ever typed by hand.

Step / Phase	Verification Focus	Passing Criteria	Result
Step 0	ADR & Spec Verification	docs/CODE_ATLAS_SPEC.md compliance	**PASS**
Step 1	Schemas & Core Contracts	Zod validation of LanguagePackage interfaces	**PASS**
Step 2	CI Pipeline & Broken Fixtures	Negative syntax traps caught deterministically	**PASS**
Step 3	App Shell & Mermaid Engine	Theme contrast AAA & SVG node generation	**PASS**
Step 4	Playground Sandbox & Watchdog	4s watchdog timeout & 64KB stdout limits	**PASS**
Step 5	Pinned Python Lessons Run	8/8 topics & 16/16 verified code examples	**PASS**
Step 6	Search v1 & Spaced Repetition	BM25 ranking & SM-2 queue calculation	**PASS**
Phase 2	Trace Studio & Grounded RAG	Temporal line frame inspection & tutor bridge	**PASS**
Phase 3	Breadth Modularity & Patterns	Multi-language comparisons & pattern templates	**PASS**
Phase 4	Accessibility & Interview Prep	Screen reader ARIA landmarks & mock algorithms	**PASS**

CI Coverage Output (`coverage/python.json`): Toolchain: Python 3.14.10 | Topics Written: 8/8 (100%) | Verified Code Examples: 16/16 (100%) | Verification Rate: 100% | CI Status: PASSED

10. Competitive Feature Matrix & Educational Impact

Compared to existing developer tools and generic chatbots, CODE ATLAS provides an unmatched combination of execution safety, visual comprehension, and cognitive learning aids:

Capability Dimension	Generic LLMs (ChatGPT / Claude)	LeetCode / HackerRank	Codecademy / Coursera	CODE ATLAS (Our Platform)
Execution Sandbox	■ Static text only; no execution	■■ Batch unit testing only; no tracing	■■ WebAssembly / browser sim only	■ **Isolated 4s watchdog sandbox**
AST Flowchart Synthesis	■ None	■ None	■ None	■ **Dynamic Mermaid topological graphs**
Error Auto-Repair	■■ Frequently hallucinates syntax	■ Opaque test failures only	■ Generic error hints	■ **Line pinpointing & 1-click verified diffs**
Temporal Line Tracing	■ None	■ None	■ None	■ **Interactive line filters & range selectors**
Spaced Repetition (SM-2)	■ None	■ None	■■ Basic streak counters	■ **Algorithmic memory decay curve modeling**
Offline Portability	■ Cloud-locked	■ Cloud-locked	■ Cloud-locked	■ **Zero-dep offline portable package (< 1MB)**

11. Instant Evaluation Guide & Future Roadmap

To ensure effortless hackathon evaluation, the entire platform has been packaged into an ultra-compact distribution archive: **CODE_ATLAS_COMPLETE.zip** (1.10 MB), containing all engine code, web assets, test runners, and visual figures.

Step 1: Extract Zip Archive: Unzip `CODE_ATLAS_COMPLETE.zip` to any local folder.

Step 2: Run CI Verification Suite: node --experimental-strip-types tools/run_complete_suite.ts (All 10 suites pass in ~8 seconds).

Step 3: Launch Local Web Server: node --experimental-strip-types tools/serve.ts (Open browser to http://localhost:3000).

Future Strategic Roadmap:

- **Q1 2027: Polyglot Micro-Adapters:** Native ExecutionAdapters and AST extractors for Rust, Go, TypeScript, and C++ with cross-language benchmarking.
- **Q2 2027: Collaborative Trace Studio:** Multiplayer peer-to-peer interactive debugging over WebRTC, allowing mentor-student live walkthroughs.
- **Q3 2027: On-Device Small Language Models (SLM):** WebLLM integration enabling 100% offline, local natural language code tutoring with zero API latency.

Conclusion: CODE ATLAS bridges the fundamental gap between static text and dynamic runtime comprehension. By enforcing grounded verification, visual topologies, and hostile code isolation, it transforms novice coders into confident, competent software engineers.

__Official Project Submission__ • CODE ATLAS Development Team • GenAI Hackathon 2026 • CODEMIND Spec v2